

SIDURI

Calibration that learns from its own alarms has no safe fixed point --- and we derive exactly how much real feedback fixes it.

$$q^*(\eta, \text{tol}) = \frac{\eta(1 - \eta - \text{tol})}{(\eta + \text{tol})(1 - \eta)}$$

Confirmed on real anomaly-detector telemetry: $\rho \approx 0.8$, $p < 10^{-6}$ (28 machines)

